Token timing under a fixed compute trace

Qwen3-0.6B | 44 prompt tokens | 8 output tokens | 10 KiB selected HBF region

(a) Modeled token-ready time

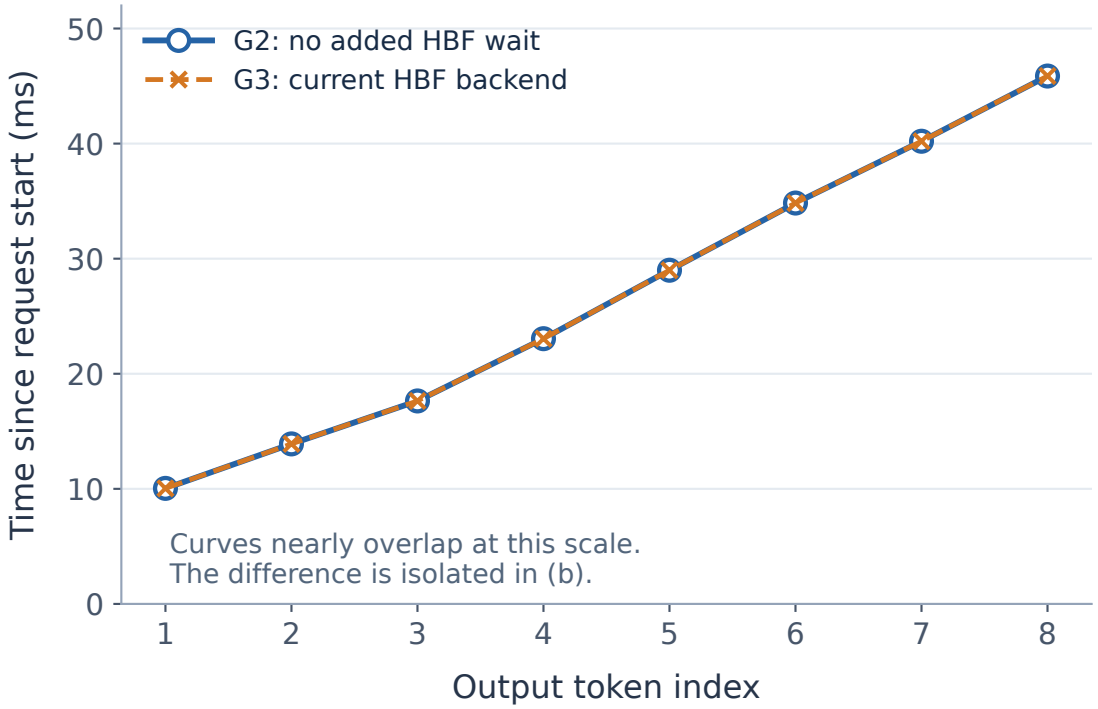

(b) Cumulative HBF increment

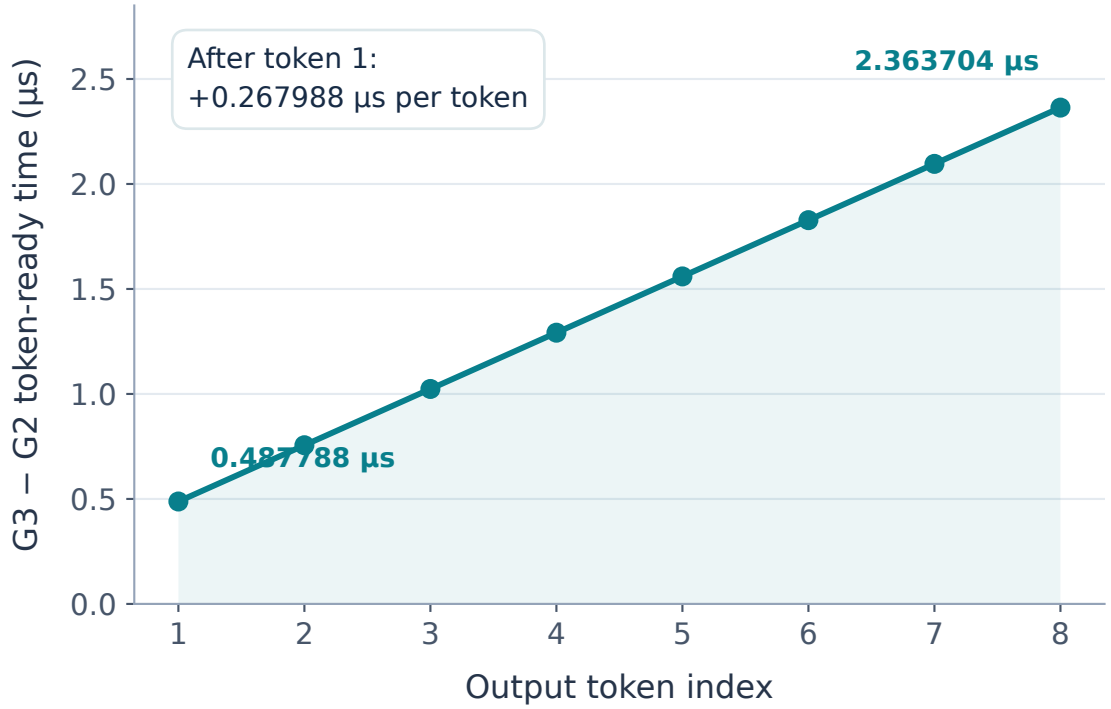

MODELED RESULTS — not measured device latency

Both cases use the same historical RTX 5070 compute table; G3 was executed on an RTX 4070 Laptop GPU.

Selected symbolic dependencies only; weights remain on GPU. H0 is uncalibrated. Lines connect token observations.